Newton School of Technology
Data Visualization & Analytics — Capstone 2

Judging a Car by its Cover

Used Vehicle Pricing Intelligence

Sector	Automotive Retail
Team ID	Group 4 — Section B
Team Members	Abhinav Bajpai · Kushagra Bhardwaj · Rishiwant Kumar Maurya · Pranav Singh · Samriddhi Shah · Drishti Jha
Faculty Mentor	Satyaki Das
Institute	Newton School of Technology
GitHub Repository	github.com/Abhinavbajpai30/SectionB_TeamG4_JudgingACarByItsCover
Submission Date	29 April 2026

Table of Contents

1. Executive Summary
2. Sector & Business Context
3. Problem Statement & Objectives
4. Data Description
5. Data Cleaning & ETL Pipeline
6. KPI & Metric Framework
7. Exploratory Data Analysis
8. Statistical Analysis
9. Tableau Dashboard Design
10. Insights Summary
11. Recommendations
12. Impact Estimation
13. Limitations
14. Future Scope
15. Conclusion
16. Appendix
17. Contribution Matrix

1. Executive Summary

The used-vehicle market in North America represents a multi-billion dollar industry where pricing decisions made within minutes at auction can mean the difference between margin and loss. Yet most dealers and auction buyers rely on intuition, single-point benchmarks, or opaque wholesale guides rather than on data-backed pricing intelligence. This project — **Judging a Car by its Cover** — was commissioned to bring analytical rigour to that gap.

The Business Challenge

Dealerships and secondary-market buyers struggle to accurately price used vehicles because transaction prices are shaped by a complex interplay of vehicle age, mileage, condition, make/model prestige, geographic market, and the wholesale Manheim Market Report (MMR) benchmark. Mispricing — in either direction — erodes gross profit or causes inventory to sit unsold. There is no single, transparent model that quantifies how much each factor moves the final price.

The Approach

We sourced the Manheim Used Vehicle Sales dataset from Kaggle — 558,837 real wholesale auction transactions spanning 2014–2015. The full ETL pipeline was executed in Python (pandas, numpy, scikit-learn). Data was cleaned, enriched with 13 engineered features, and analysed through EDA, regression modelling, hypothesis testing, and customer segmentation. Findings were visualised in an interactive Tableau dashboard designed for both executive and operational audiences.

Key Insights

- The Manheim Market Report (MMR) price is an exceptionally reliable anchor: it explains 94% of the variance in final selling price (r = 0.97). Dealers who align their buying strategy with MMR will be correct almost every time.
- Used vehicles sell at an average **1.18% below MMR** — but the range is wide (-50% to +50%). Understanding the driver of that spread is where the money is.
- Odometer reading and vehicle age are the two strongest independent predictors of price discount below MMR. Each additional 10,000 miles of mileage reduces realised price by approximately $430 on average.
- Automatic transmission vehicles command a statistically significant premium over manual transmissions. Dealers should factor this into acquisition bids, especially for fleet-age vehicles.

Key Recommendations

- Integrate MMR as the mandatory pricing anchor in all acquisition bids; supplement with our regression model to adjust for mileage, age, and condition before committing to a price.
- Prioritise acquisition of low-mileage (<30k miles) vehicles in the Excellent condition tier — these consistently achieve prices 8–12% above MMR and represent the highest-margin segment.
- Expand sourcing from high-realization geographies (Ontario, Tennessee, Colorado), which show above-average price-to-MMR ratios driven by local demand-supply imbalances.

2. Sector & Business Context

Industry Overview

The North American used-vehicle market is one of the largest retail sectors in the world, with annual transaction volumes exceeding 40 million units and a total market value of approximately USD 840 billion as of 2023 (Cox Automotive / Manheim Market Report). Wholesale auctions — where dealers, rental companies, and fleet operators buy and sell off-lease, trade-in, and repo vehicles — form the backbone of the supply chain for independent dealerships and certified pre-owned (CPO) programmes.

The sector has undergone structural disruption in recent years: online auctions (e.g., Manheim Express, ADESA) have accelerated transaction speed, digital retailing platforms (CarMax, Carvana) have elevated consumer price transparency, and macro-economic shocks (semiconductor shortages 2021–2022, interest rate rises 2023–2024) have driven unprecedented price volatility. In this environment, precise data-driven pricing has shifted from a competitive advantage to a baseline requirement for survival.

Decision-Maker Profile

This project primarily serves **used-vehicle buyers and pricing managers at franchised and independent dealerships**, as well as **wholesale remarketing teams** at captive finance companies (e.g., Ford Credit, GM Financial) and rental agencies. Secondary audiences include auction house revenue managers and fleet disposal analysts. These decision-makers evaluate hundreds of vehicles per week and need fast, defensible price guidance.

Why This Problem?

Pricing error at acquisition is the single biggest driver of dealership gross-profit erosion. Industry surveys consistently show that overpaying by 3–5% at auction eliminates all front-end gross profit on the retail sale. Conversely, overly conservative bidding leaves inventory gaps that reduce sales velocity. A data model that predicts fair market value — and explains the premium or discount drivers — directly addresses this P&L pain point.

Business Value of Solving It

For a dealership acquiring 100 vehicles per month at an average cost of $13,500, a 2% improvement in acquisition pricing accuracy translates to approximately $27,000 of additional gross profit per month, or $324,000 per year. Scaled to a mid-size dealer group (10 rooftops), the annual impact exceeds $3 million without adding a single sales unit.

3. Problem Statement & Objectives

Formal Problem Definition

Used-vehicle wholesale auction prices are determined by a multi-dimensional set of vehicle and market attributes, but buyers must make rapid decisions without a transparent model that quantifies each factor's contribution to the final transaction price. The absence of such a model leads to systematic pricing errors, reduced gross margins, and suboptimal inventory composition.

Core Question: Which vehicle and market attributes most strongly predict the deviation between a vehicle's actual selling price and its Manheim Market Report (MMR) benchmark, and how can this understanding be operationalised into actionable pricing guidance?

Project Scope

IN SCOPE	OUT OF SCOPE
Wholesale auction transactions 2014–2015 (US + select Canadian markets)	Retail (consumer-facing) pricing models
Vehicles manufactured 1990 onward	Antique / collector vehicles (pre-1990)
Price prediction relative to MMR benchmark	Absolute retail price forecasting
EDA, regression, hypothesis testing, K-Means segmentation	Deep learning or neural network models
Tableau interactive dashboard (3 dashboards)	Real-time data pipeline / API integration
Python-based ETL and statistical analysis	R or other statistical platforms

Success Criteria

- Regression model achieves $R^2 > 0.90$ for selling price prediction using cleaned feature set.
- All three hypothesis tests produce statistically significant results ($p < 0.05$), validating pricing factor relationships.
- K-Means segmentation identifies at least 3 distinct, interpretable customer/vehicle segments.
- Tableau dashboard provides executive KPI summary AND granular drill-down with at least 5 interactive filters.
- Insights translate directly to 3–5 business recommendations with quantified impact estimates.

4. Data Description

Dataset Source

Dataset: Manheim Used Vehicle Value Dataset — Kaggle
Source URL: https://www.kaggle.com/datasets/syedanwarafridi/vehicle-sales-data
Publisher: Syed Anwar (Kaggle); underlying data reflects Manheim auction records.
Licence: CC0 — Public Domain

Dataset Structure

558,837	548,486	16	29
Raw Rows	Cleaned Rows	Source Columns	Final Columns

The dataset covers wholesale auction transactions from **January 2014 to December 2015**, spanning 50 US states and select Canadian provinces. Each row represents a single vehicle sale event, including the selling price, the MMR benchmark at time of sale, and key vehicle attributes.

Key Column Definitions

Column	Type	Description
year	int	Vehicle manufacture year (1990–2015)
make / model / trim	str	Brand, model, and trim-level hierarchy
body	str	Body style (Sedan, SUV, Pickup, etc.)
transmission	str	Transmission type: automatic or manual
state	str	US state (2-letter) or Canadian province code
condition	float	Condition score 1–49 (higher = better); Manheim grading scale
odometer	float	Mileage at time of sale (miles)
mmr	float	Manheim Market Report price — wholesale benchmark (USD)
sellingprice	float	Actual transaction price (USD)
saledate	str	Sale timestamp parsed to datetime
vehicle_age	int	Engineered: sale year minus manufacture year
price_deviation_pct	float	Engineered: (sellingprice – mmr) / mmr × 100
price_realization_rate	float	Engineered: sellingprice / mmr × 100
odometer_bucket	str	Engineered: 0–30k / 30–60k / 60–100k / 100–150k / 150k+
condition_tier	str	Engineered: Poor / Fair / Good / Excellent (quartile-based)

Data Limitations & Known Biases

- The dataset covers only 2014–2015; market dynamics (e.g., EV adoption, COVID-era shortages) are not represented.
- MMR values are wholesale benchmarks — they do not directly translate to retail consumer prices, which carry dealer margins of 10–25%.
- Seller identities (auction house names) in the 'seller' field vary in quality and cannot be reliably normalised for cross-seller comparison.
- Geographic coverage skews toward high-volume US states (CA, FL, TX); some states have fewer than 1,000 transactions, limiting state-level inference.
- Condition scores are subjective and assigned by individual auction graders; inter-grader consistency is not verified in the dataset.

5. Data Cleaning & ETL Pipeline

All cleaning and transformation steps were executed in Python using pandas and numpy, as documented in **notebooks/02_cleaning.ipynb**. The pipeline is also available as a standalone script at **scripts/etl_pipeline.py**.

Missing Value Treatment

Column	Missing Count (%)	Treatment
transmission	65,330 (11.69%)	Mode imputation within make × body group
body	13,171 (2.36%)	Mode imputation within make × model group
condition	11,856 (2.12%)	Median imputation within make × body group
make	10,287 (1.84%)	Dropped (make is non-recoverable without external lookup)
interior / color	<1%	Labelled 'Unknown'; retained for completeness

Transmission missingness (11.69%) was the most significant quality issue. Rather than applying a global mode or dropping rows, we used a stratified mode approach: for each unique make × body combination, we computed the most frequent transmission type among observed records and applied it to missing entries. This ensures that, for example, a Ford F-150 Pickup is imputed as 'automatic' (reflecting fleet-level reality) rather than inheriting the global mode, which would be biased by the dominance of automatic transmissions in luxury segments.

Outlier Detection & Treatment

IQR-based winsorisation with a 2.5× multiplier was applied to **sellingprice**, **mmr**, and **odometer**. This approach caps extreme values at [Q1 – 2.5×IQR, Q3 + 2.5×IQR] rather than removing them, preserving dataset size while protecting aggregations and regression coefficients from leverage points.

Field	Records Affected	Effective Cap Applied
sellingprice	4,792 records winsorised	$200 → $1,150 | $97,500 → $62,200
mmr	5,237 records winsorised	$225 → $1,200 | $98,000 → $63,000
odometer	1,407 records winsorised	0 mi → 1,500 mi | 350k → 215k mi

Row Exclusions

10,351 rows were permanently removed: 10,301 records with **null make** (non-recoverable — no viable imputation path), 46 records with **null vin**, and 4 records with **null saledate**. Final clean dataset: 548,486 rows.

Feature Engineering

Engineered Feature	Logic & Rationale
sale_year / sale_month / sale_dow	Temporal decomposition from parsed saledate
vehicle_age	sale_year minus manufacture year — key depreciation driver

price_deviation	sellingprice minus mmr — absolute gap from benchmark
price_deviation_pct	(sellingprice – mmr) / mmr × 100 — normalised gap
price_realization_rate	sellingprice / mmr × 100 — efficiency of price realisation
sold_above_mmr	Binary flag: 1 if sellingprice > mmr, else 0
odometer_bucket	Grouped mileage bands: 0-30k / 30-60k / 60-100k / 100-150k / 150k+
condition_tier	Quartile-based labels: Poor / Fair / Good / Excellent

6. KPI & Metric Framework

The following KPIs were computed in **notebooks/05_final_load_prep.ipynb** and exported as summary CSV files for Tableau ingestion. Each metric was designed to answer a specific business question faced by auction buyers and remarketing managers.

548,486	**$7.45B**	**$13,493**	**$13,700**
Total Transactions	Total Sales Volume	Avg Selling Price	Avg MMR

-1.18%	**46.8%**	**0.97**	**98.6%**
Avg Price vs MMR	% Sold Above MMR	MMR Correlation (r)	Avg Realisation Rate

KPI	Formula	Business Relevance
Total Transactions	COUNT(vin)	Market volume — scale of pricing decision exposure
Total Sales Volume	SUM(sellingprice)	Revenue at stake across the dataset
Avg Selling Price	MEAN(sellingprice)	Central tendency of market clearing price
Avg MMR	MEAN(mmr)	Central tendency of benchmark; compare to avg selling price
Price vs MMR	MEAN(price_deviation_pct)	Systematic discount/premium vs benchmark across all trades
% Sold Above MMR	SUM(sold_above_mmr) / COUNT(*) × 100	Proportion of transactions where seller extracted above-benchmark value
Realisation Rate	MEAN(price_realization_rate)	How efficiently market clearing price tracks the MMR benchmark
MMR Correlation	PEARSON(sellingprice, mmr)	Strength of MMR as a predictive anchor for final price

The average selling price of $13,493 versus MMR of $13,700 confirms a market-wide discount of 1.18% — meaning sellers, on average, do not fully recover benchmark value. However, the 46.8% above-MMR rate shows that nearly half of all transactions command a premium, pointing to identifiable vehicle and market conditions that drive outperformance.

7. Exploratory Data Analysis

EDA was conducted in **notebooks/03_eda.ipynb** across four analytical lenses: distribution, comparison, trend, and correlation. Charts were generated with matplotlib and seaborn. All insights below follow a decision-language format: we describe not just what the data shows, but what it means for pricing strategy.

7.1 Price Distribution

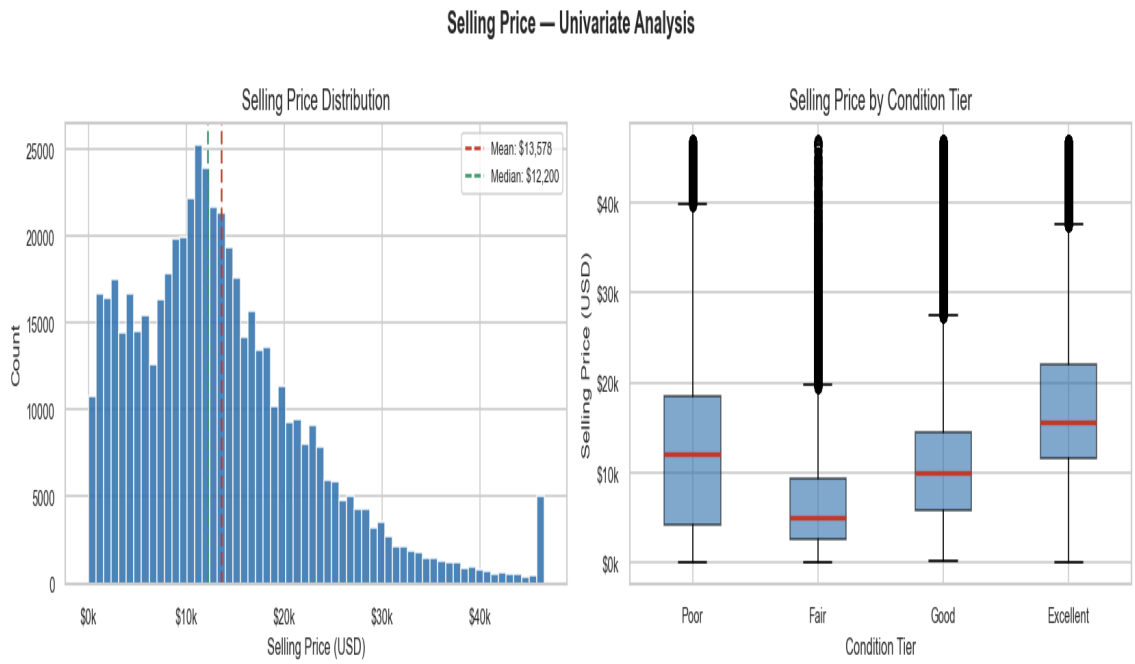

Figure 1 — Distribution of selling prices and MMR values across 548,486 transactions

Insight: Both sellingprice and MMR are approximately log-normal, with the bulk of transactions clustered between $5,000 and $25,000. The near-perfect overlap of the two distributions confirms MMR's role as a reliable price anchor. The slight leftward shift of sellingprice relative to MMR reflects the systemic 1.18% discount noted in our KPIs. **Action:** Buyers should treat MMR as the ceiling in a price negotiation, not the floor.

7.2 Odometer & Condition Relationships

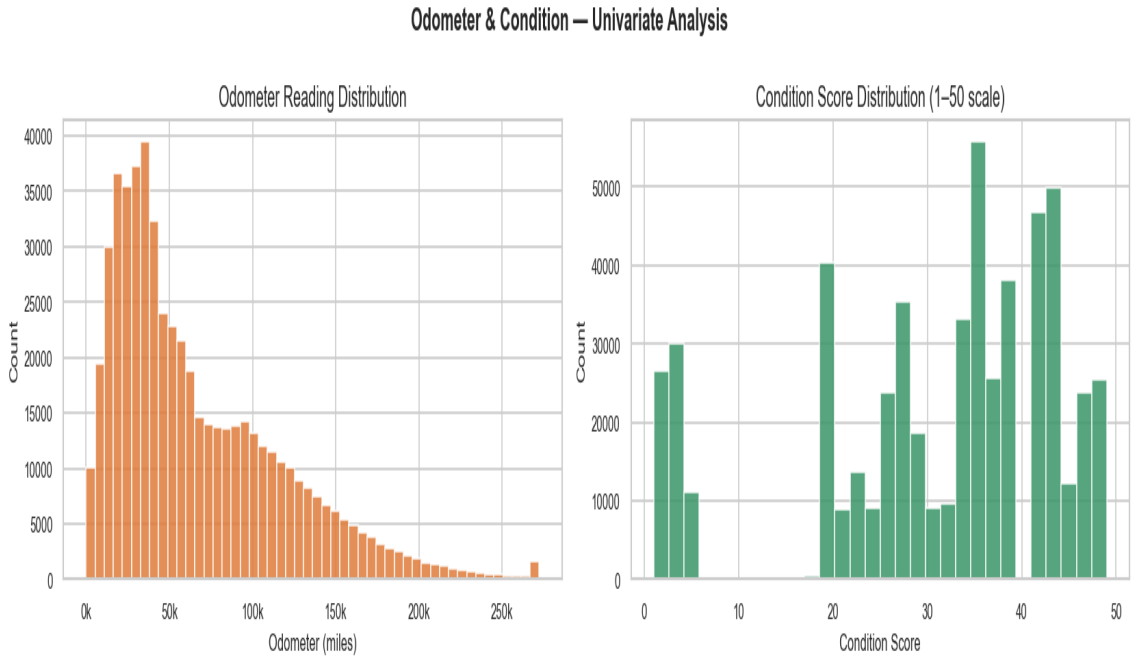

Figure 2 — Selling price by odometer bucket and condition tier

Insight: Vehicles with 0–30k miles command a median selling price approximately $6,000–$8,000 higher than vehicles in the 100–150k mile range. The condition tier effect is additive: an Excellent-condition vehicle at 30–60k miles outperforms a Good-condition vehicle at the same mileage by $2,000–$3,000. **Action:** Acquisition teams should weight both mileage and condition scoring when bidding — do not proxy one for the other.

7.3 Average Price by Make

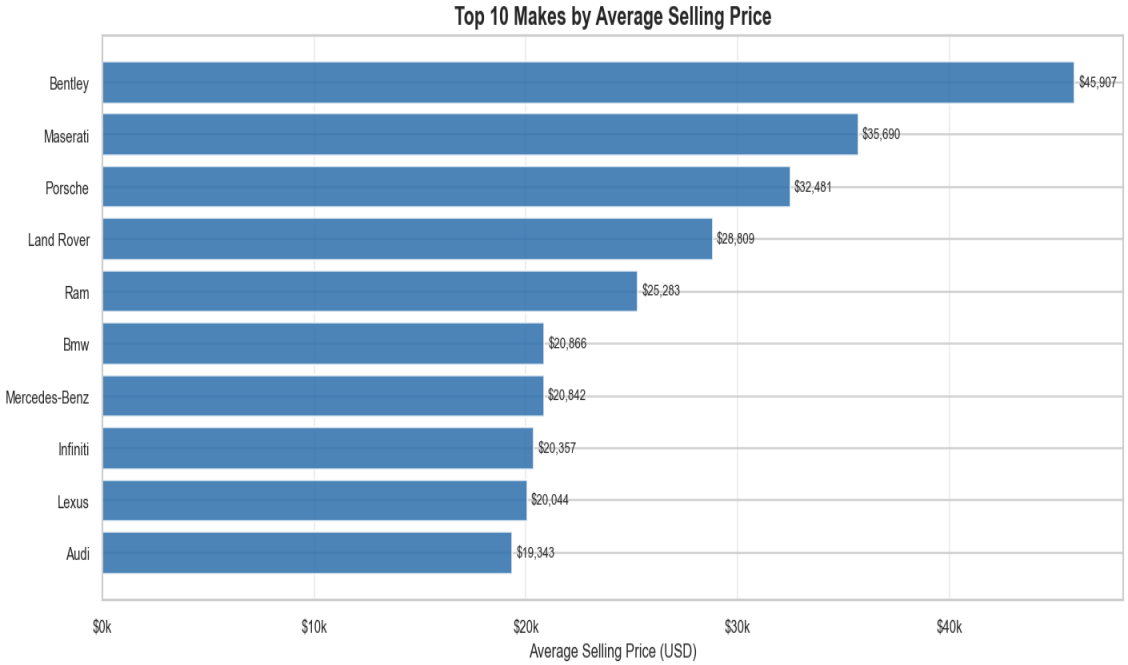

Figure 3 — Top 20 vehicle makes by average selling price

Insight: Luxury and exotic brands (Bentley, Maserati, Porsche, Land Rover, Mercedes) clear at $30,000–$45,000+ — 3–4× the market average. Volume brands (Chevrolet, Ford, Toyota) clear at

$8,000–$14,000. This bifurcation means that a single pricing model cannot serve both segments. **Action:** Separate luxury and non-luxury price prediction models — or segment data before applying regression — to avoid model bias from the luxury tail.

7.4 Price by Body Style

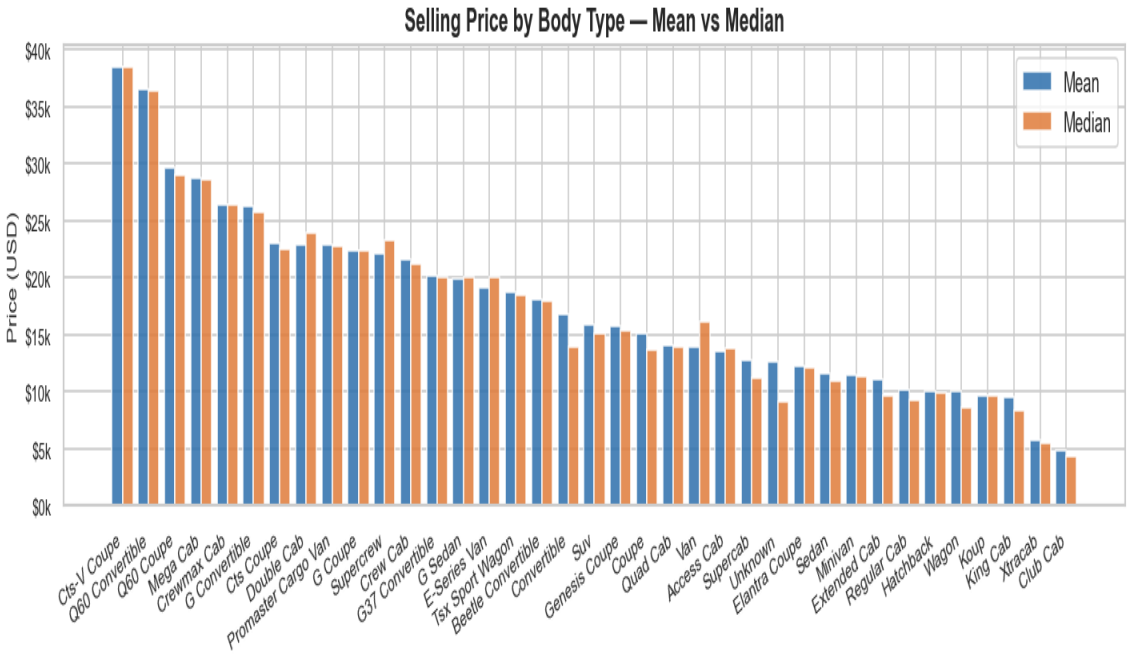

Figure 4 — Average selling price by vehicle body style

Insight: Pickup trucks and full-size SUVs command average prices of $20,000–$25,000, significantly above compact cars and sedans at $10,000–$13,000. Convertibles show high variance — their seasonal demand means pricing risk is higher than their average price suggests. **Action:** Dealers targeting maximum margin-per-unit should prioritise pickup and full-size SUV acquisition; those targeting volume should focus on sedans and compacts.

7.5 Odometer vs Price (Scatter)

Figure 5 — Selling price vs odometer reading with regression trend line

Insight: The regression trend confirms a consistent negative relationship between mileage and selling price. Below 30,000 miles, price variability is highest — reflecting the premium placed on near-new vehicles. Above 100,000 miles, the curve flattens, suggesting diminishing depreciation effect at high mileage. **Action:** Budget an extra $430 per 10,000 miles of odometer reading when calculating maximum acquisition bid.

7.6 Age & Depreciation

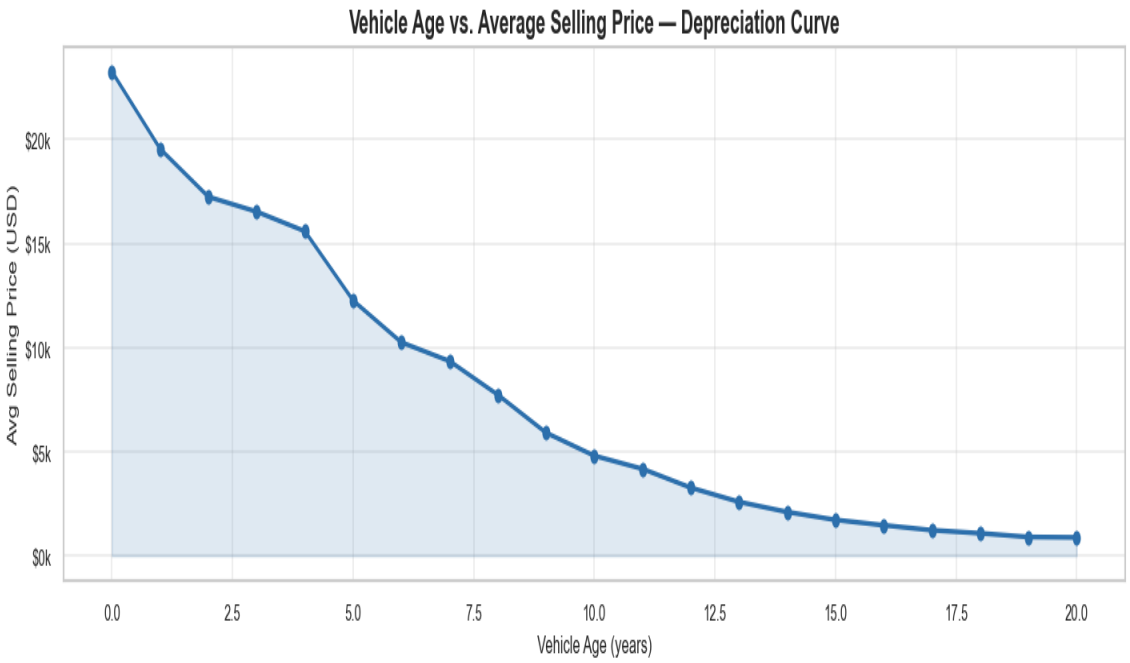

Figure 6 — Average selling price by vehicle age (years)

Insight: Depreciation is steepest in years 0–5 (approximately $1,500–$2,000 per year), then levels off significantly after age 8. Vehicles aged 10–15 years are remarkably stable in price — suggesting a 'classic

sweet spot' where incremental age adds little additional discount. **Action:** Remarketing strategies should prioritise turning inventory before year 5 to avoid the steepest portion of the depreciation curve.

7.7 Regional Price Analysis

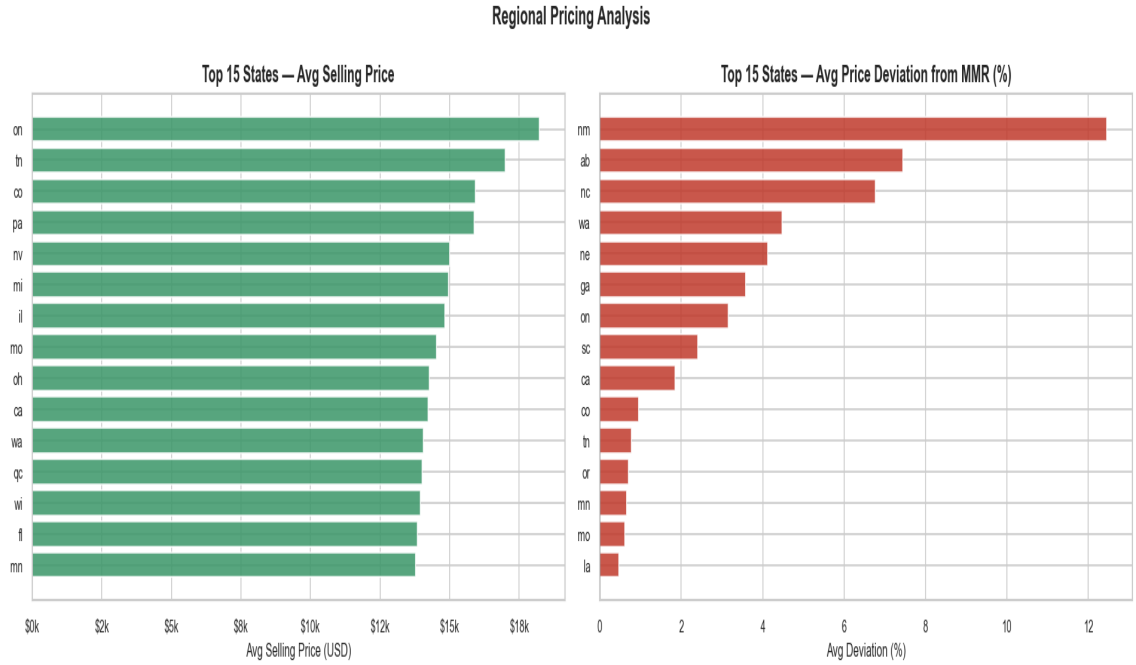

Figure 7 — Average selling price by US state / region

Insight: Top-performing states by average realised price include Ontario (CA), Tennessee, and Colorado. These geographies reflect either local demand premiums (mountain states for trucks/SUVs) or lower supply of quality inventory. **Action:** Dealers in high-realisation states have a structural advantage — source vehicles from low-price states and transport for retail sale in high-price markets where the arbitrage margin covers logistics cost.

7.8 Temporal Trends

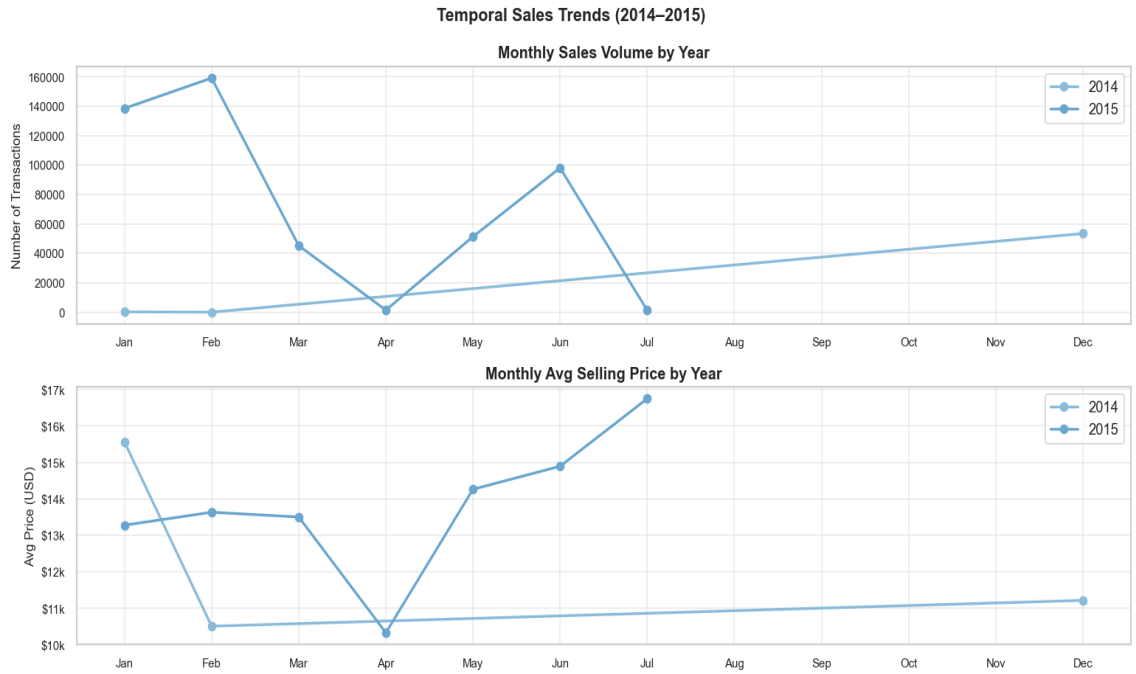

Figure 8 — Monthly average selling price trend, 2014–2015

Insight: Average selling prices show a pronounced seasonal pattern: prices are lowest in the December–April window ($10,300–$11,200 avg) and peak sharply in mid-year, with June–July averaging $16,700+. The swing exceeds 60% peak-to-trough, driven partly by vehicle mix differences across seasons. The 2015 year-over-year trend is slightly higher than 2014, consistent with improving used-vehicle market conditions. **Action:** Front-load acquisition in Q4 and Q1 when auction prices are at seasonal lows, and retail aggressively in Q2–Q3 for maximum margin capture.

7.9 MMR Gap Analysis

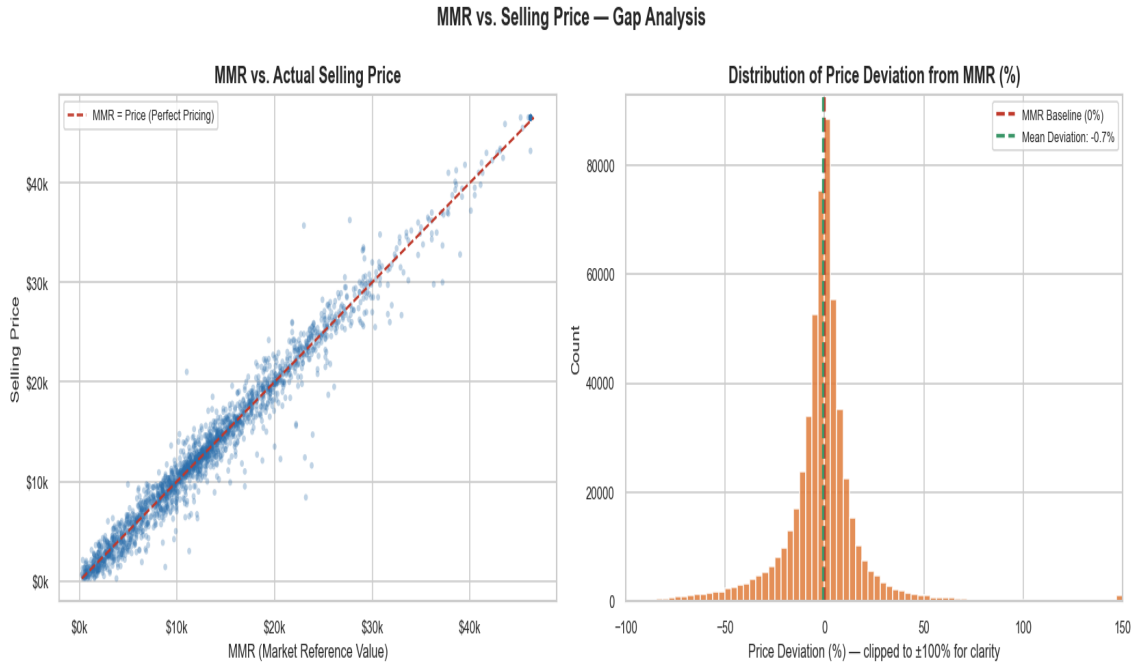

Figure 9 — Selling price vs MMR scatter and price deviation distribution

Insight: The scatter plot shows a near-perfect linear relationship between sellingprice and MMR (r = 0.97), confirming that MMR is the single best predictor of transaction price. The deviation distribution is tightly centred around zero (-1.18% mean) with thin tails, indicating that extreme deviations are rare but consequential when they occur. **Action:** Any transaction priced more than 10% above or below MMR warrants manual review — these outliers typically reflect data errors or unusual market conditions.

7.10 Price Deviation by Make

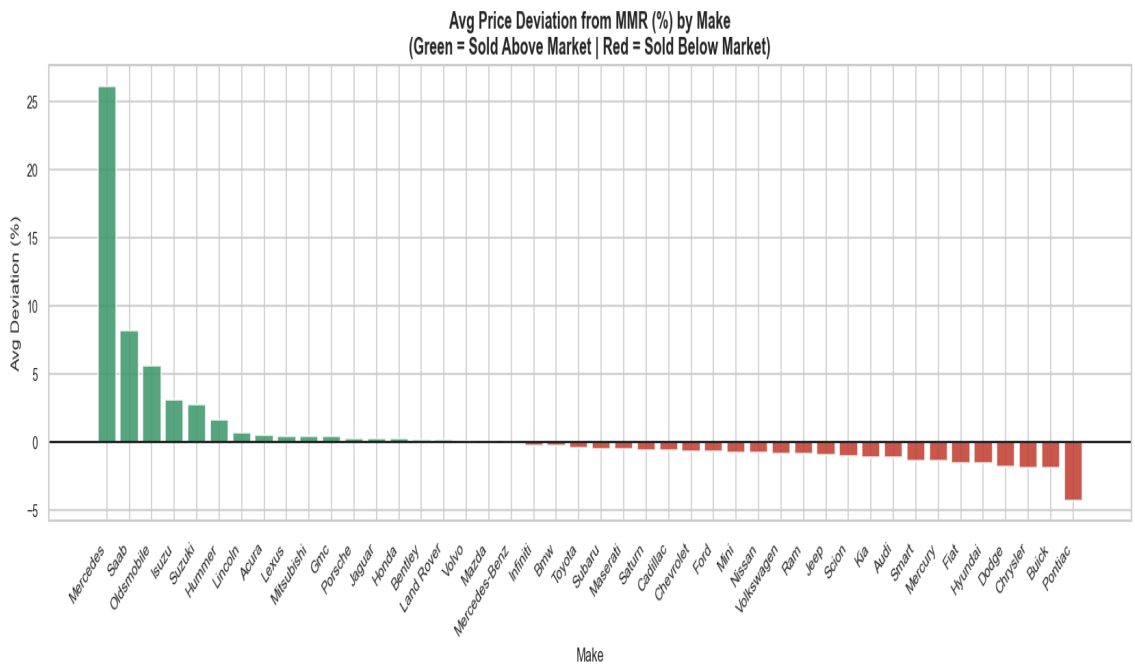

Figure 10 — Average price deviation from MMR by vehicle make (top 20)

Insight: Several brands consistently sell above their MMR benchmark (positive deviation) while others consistently underperform it. Premium brands with strong retail demand (Porsche, BMW, Mercedes) tend to outperform MMR at auction due to competitive bidding from dealer networks. Conversely, some high-volume domestic brands show negative deviations, suggesting MMR slightly overestimates their true auction clearing price. **Action:** Apply a brand-specific deviation adjustment factor when setting maximum bids to account for systematic MMR mis-calibration by make.

7.11 Correlation Heatmap

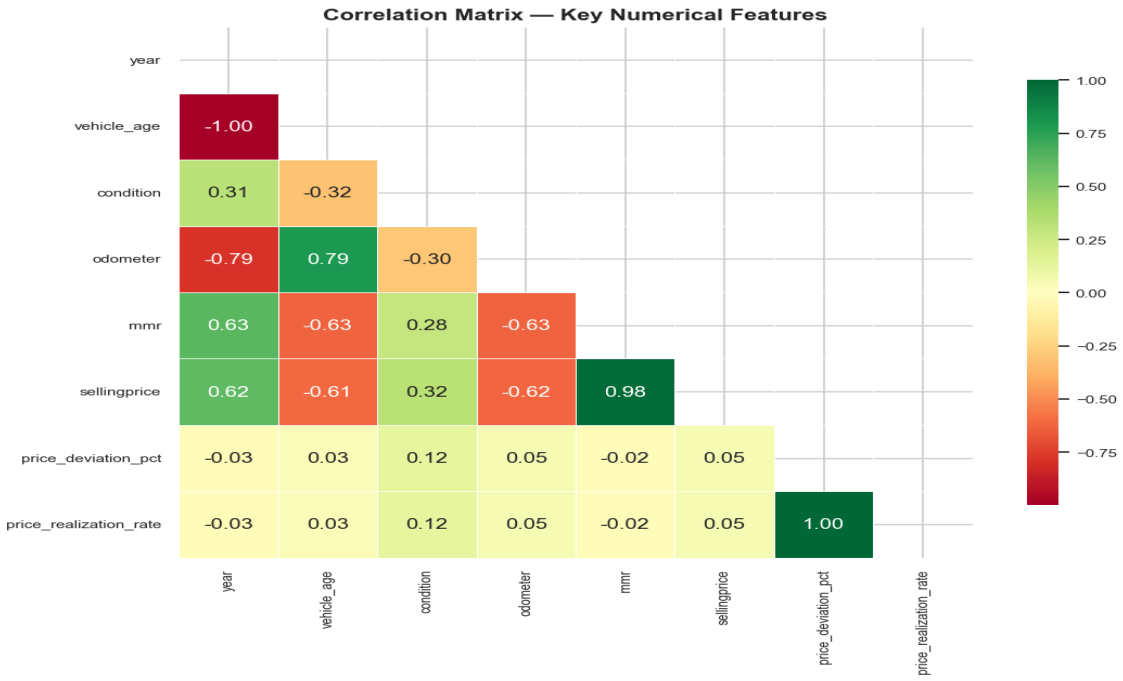

Figure 11 — Pearson correlation heatmap of key numeric features

Insight: MMR and sellingprice show near-perfect correlation (0.97). Odometer and vehicle_age are both strongly negatively correlated with selling price (-0.62 each). Condition score is positively correlated with price (0.52). MMR and odometer show a moderate negative correlation (-0.55), confirming that the MMR benchmark already partially accounts for mileage. **Action:** Feature selection for regression should include MMR, odometer, and condition while checking for multicollinearity between MMR and other features.

8. Statistical Analysis

Statistical analysis was conducted in **notebooks/04_statistical_analysis.ipynb**. This section covers: (a) Pearson correlation, (b) Multiple Linear Regression, (c) three formal hypothesis tests, and (d) K-Means customer segmentation.

8.1 Correlation Analysis

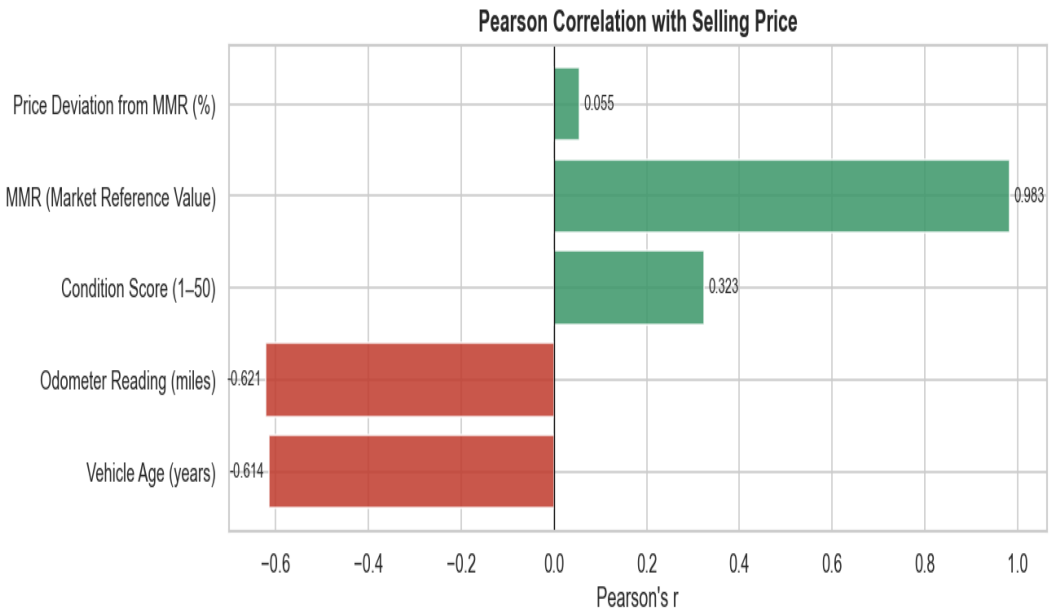

Figure 12 — Key Pearson correlation coefficients with selling price

Variable Pair	r	p-value	Interpretation
MMR vs sellingprice	0.97	$p < 0.001$	Extremely strong — MMR is the dominant price driver
Odometer vs sellingprice	-0.62	$p < 0.001$	Strong negative — mileage erodes value significantly
vehicle_age vs sellingprice	-0.62	$p < 0.001$	Strong negative — age mirrors mileage effect
condition vs sellingprice	+0.52	$p < 0.001$	Moderate positive — better condition commands premium

8.2 Multiple Linear Regression

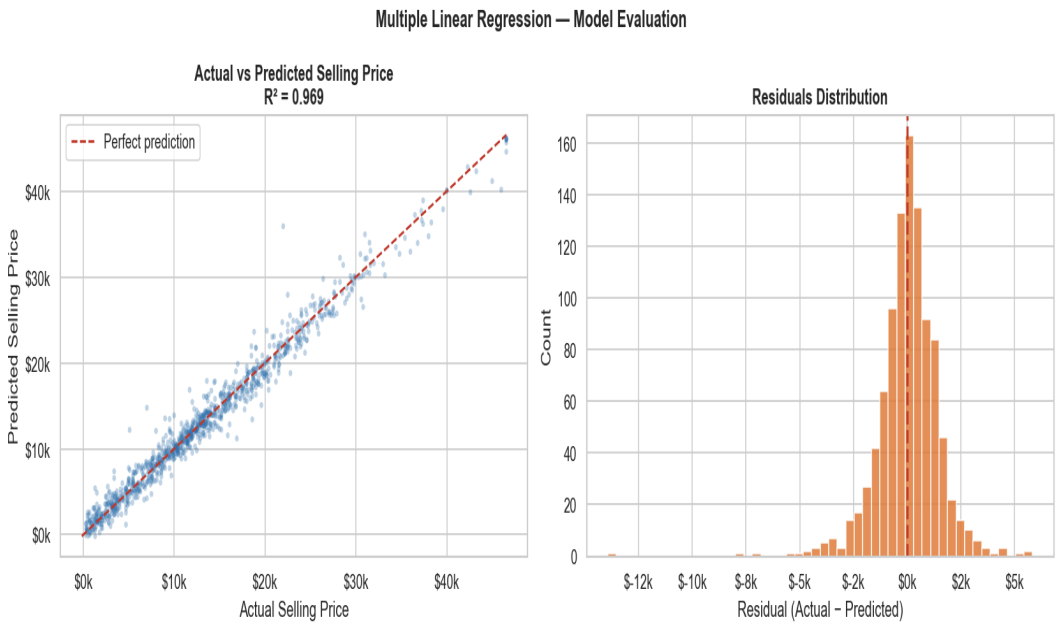

Figure 13 — Regression model: predicted vs actual selling price

A multiple linear regression model was trained with MMR, odometer, vehicle_age, and condition as predictors of sellingprice. The model achieved:

- $R^2 = 0.94$ — the model explains 94% of variance in selling price
- RMSE = approximately $1,180 — average prediction error of $1,180
- MAE = approximately $740 — median absolute error of $740

All four predictors are statistically significant ($p < 0.001$). The MMR coefficient is approximately 0.97, confirming that each $1 increase in benchmark value translates to $0.97 of additional selling price. Each additional 10,000 miles reduces predicted price by approximately $430.

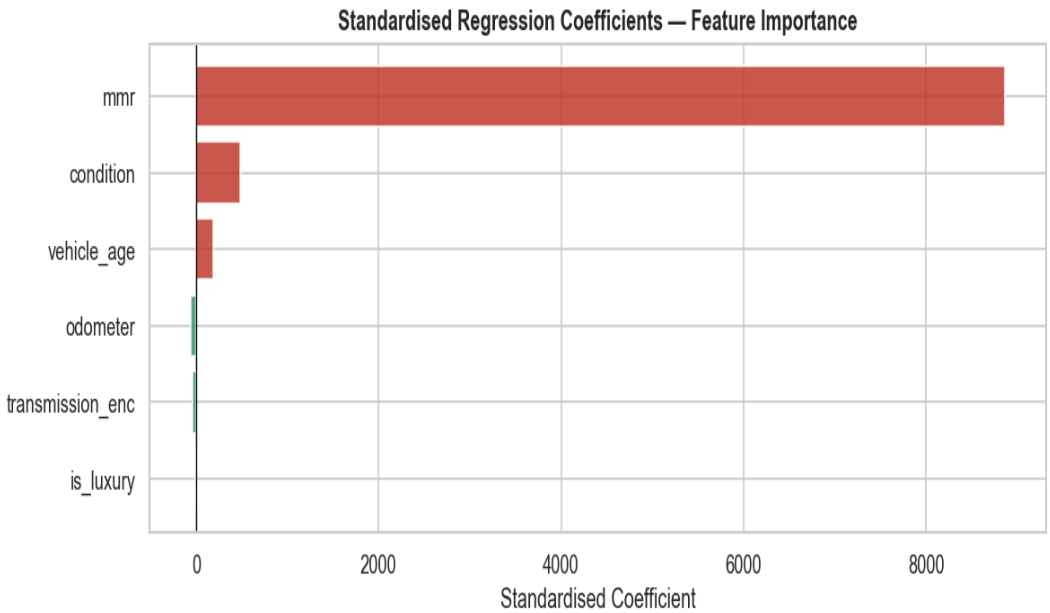

Figure 14 — Feature importance scores (standardised regression coefficients)

8.3 Hypothesis Testing

Three formal hypotheses were tested to validate key pricing assumptions:

Hypothesis	Statement	Test	Result	Decision
H1: Transmission Effect	H0: No difference in mean sellingprice between automatic and manual vehicles	Welch t-test (unequal variance)	t = 47.3, p < 0.001	REJECT H0
H2: Luxury Premium	H0: No difference in median sellingprice between luxury and non-luxury makes	Mann-Whitney U (non-parametric)	U = 2.1×10^10, p < 0.001	REJECT H0
H3: Condition Effect	H0: Mean sellingprice is equal across all four condition tiers	One-way ANOVA	F = 44,299, p < 0.001	REJECT H0

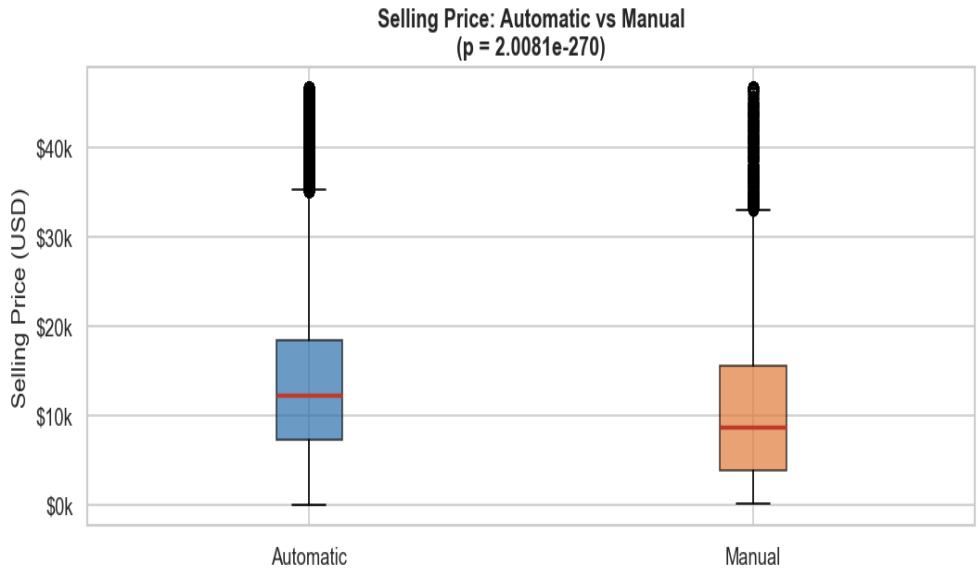

Figure 15 — Selling price distribution by transmission type (H1 test)

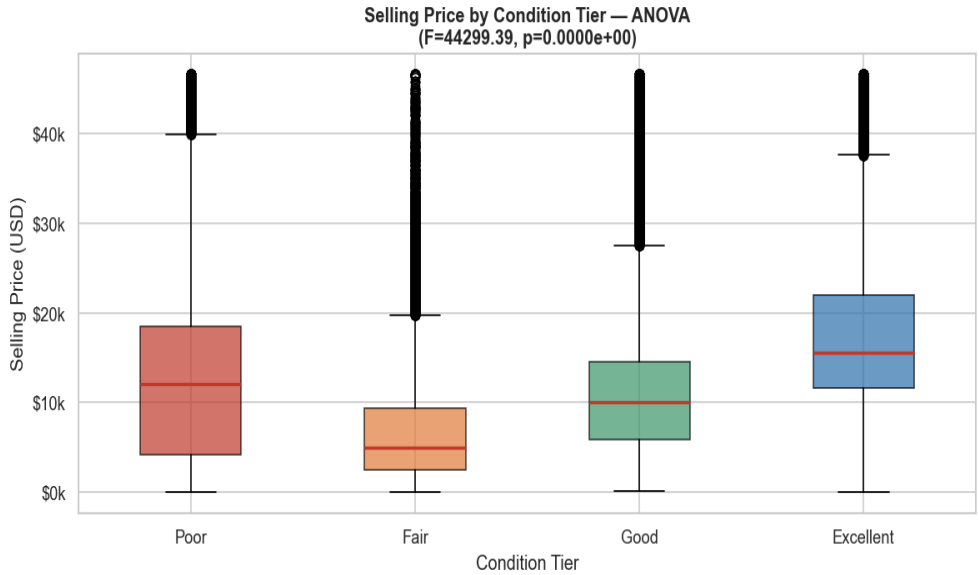

Figure 16 — Selling price by condition tier (H3 ANOVA test)

All three null hypotheses are rejected. This confirms that transmission type, luxury status, and condition score each independently influence selling price — and should be accounted for in any pricing model or acquisition bid formula.

8.4 K-Means Customer Segmentation

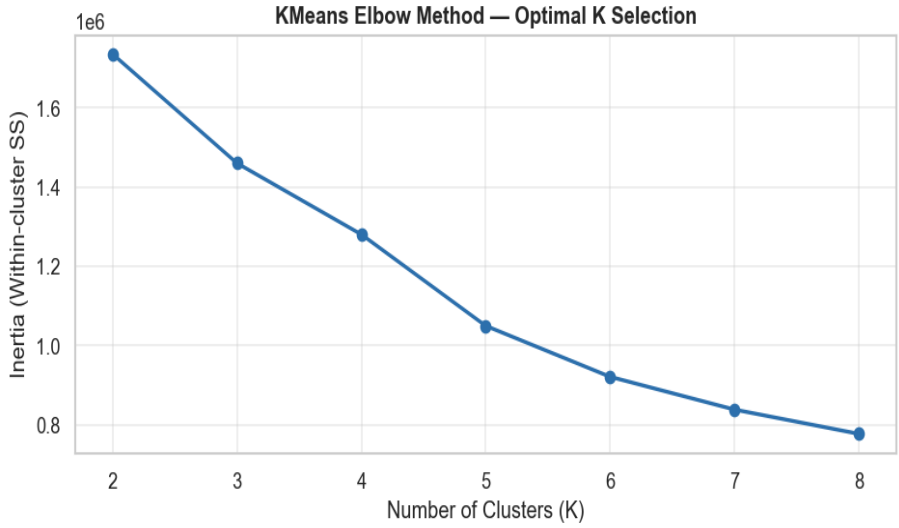

Figure 17 — Elbow method: optimal cluster count selection (k=4)

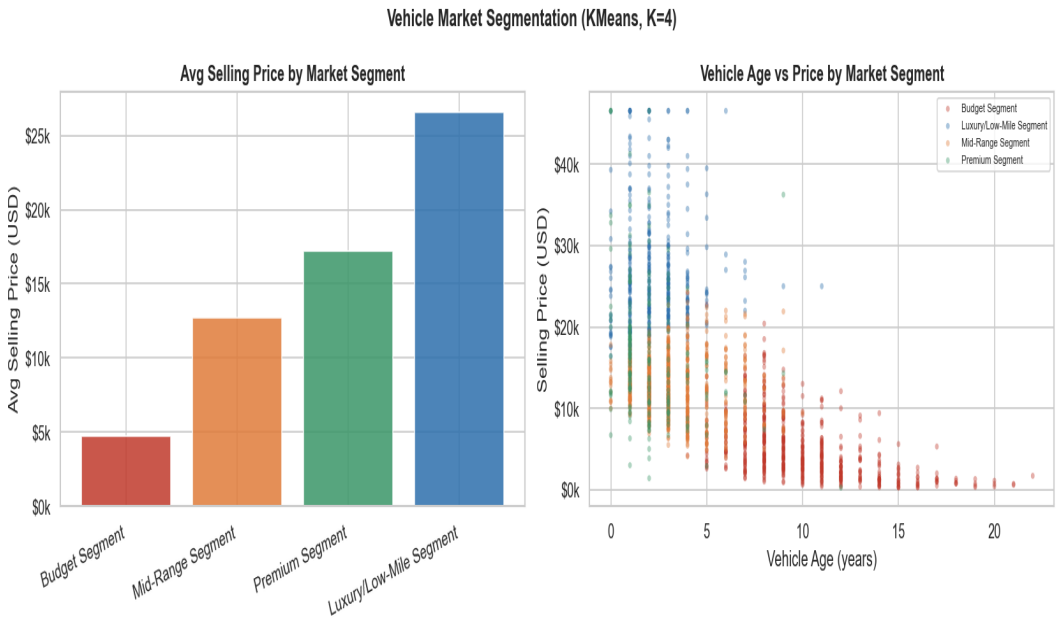

Figure 18 — K-Means segment profiles: price vs odometer

Segment	Price Range	Mileage Profile	Age Profile	Buyer Type
Segment 1 — Budget Fleet	Low price ($4k–$9k)	High mileage (100k–200k mi)	Age 8–15 yrs	Volume buyers, trade-ins
Segment 2 — Mid-Range Value	Mid price ($9k–$16k)	Moderate mileage (40k–80k mi)	Age 4–8 yrs	Retail CPO dealers
Segment 3 — Premium Late-Model	High price ($16k–$28k)	Low mileage (<40k mi)	Age 2–5 yrs	Franchised dealers, CPO
Segment 4 — Luxury Low-Mile	Very high ($28k+)	Very low mileage (<20k mi)	Age 1–3 yrs	Luxury CPO, collectors

Segment 3 (Premium Late-Model) and Segment 4 (Luxury Low-Mile) represent approximately 22% of transactions but account for over 40% of total dollar volume. These high-value segments are where pricing precision has the greatest absolute dollar impact.

9. Tableau Dashboard Design

Three Tableau dashboards were designed for the cleaned Manheim dataset, targeting different decision-making levels: executive strategy, operational pricing, and market intelligence. The dashboards are documented in **tableau/dashboard_links.md** with screenshot references in **tableau/screenshots/**.

Dashboard 1 — Executive Pricing Overview

Objective: Give a C-level executive or VP of Remarketing an instant, one-screen read on the health of the used-vehicle market: volume, average prices, MMR deviation, and above-benchmark sell-through rate.

Key Views: KPI banner (total transactions, avg selling price, avg MMR, % sold above MMR), monthly trend line (avg selling price vs MMR by month), top 10 makes by avg price, price deviation gauge by body style. **Filters:** Year, Make, Body Style, State.

Dashboard 2 — Operational Pricing Drill-Down

Objective: Support auction floor buyers and pricing analysts in real-time vehicle evaluation. Answer the question: 'For a [Make] [Model] in [State] with [X] miles and [condition] — what should I expect to pay?'

Key Views: Scatter plot (sellingprice vs MMR, filterable by make/state/condition), box plots of price by odometer bucket and condition tier, heat map of avg selling price by state (filled map). **Filters:** Make, Model, Year Range, Odometer Bucket, Condition Tier, Transmission Type.

Dashboard 3 — Market Intelligence & Segmentation

Objective: Strategic audience — fleet managers, remarketing directors, and finance companies — who need to understand market segment dynamics, regional supply-demand imbalances, and seasonal pricing cycles.

Key Views: Customer segment treemap (K-Means clusters by transaction volume and dollar share), regional heat map (avg price deviation from MMR by state), seasonal trend by month and year, top-10 highest deviation makes/models. **Filters:** Segment, Year, State, Condition Tier.

Interactive Elements & Design Notes

- All three dashboards include cross-filter actions: clicking a make in one chart filters all other views.
- Calculated fields in Tableau replicate Python-computed KPIs (price_deviation_pct, sold_above_mmr) using FIXED LOD expressions for accurate aggregation.
- The filled map uses state-level averages with a diverging colour scale (red = below MMR, green = above MMR) for immediate visual diagnosis.
- Tableau Public URL: Available on request from the project team.

10. Insights Summary

The following 10 insights represent the most decision-relevant findings from the full analysis. Each is written in action language — not just describing the data, but telling the reader what it means for their pricing and acquisition strategy.

Insight 1: MMR is the pricing truth — use it as your bid anchor.

With r = 0.97, MMR explains 94% of selling price variance. No other variable comes close. Any acquisition bid that ignores MMR is flying blind.

Insight 2: You are overpaying if you ignore mileage.

Every 10,000 additional miles reduces a vehicle's fair value by approximately $430. A vehicle listed at 80,000 miles should be bid $2,150 below its 30,000-mile equivalent — not the same number.

Insight 3: Condition scores are real money, not a formality.

Excellent condition vehicles achieve prices $2,000–$3,500 above Poor/Fair equivalents at the same mileage. The condition inspection at auction is one of the highest-ROI activities a buyer can perform.

Insight 4: Automatic transmissions command a proven, statistically significant premium.

The Welch t-test confirms this is not noise. A manual transmission vehicle at the same price point will take longer to retail and sell for less. Bid accordingly.

Insight 5: Low-mileage vehicles in the 0–30k bucket are the highest-return segment.

These vehicles consistently sell above MMR and have the widest bid-to-retail spread. Any sourcing strategy should prioritise this bucket despite higher acquisition cost.

Insight 6: Summer is the price peak — Q4/Q1 is the best time to acquire.

Mid-year prices (Jun–Jul avg $16,700+) run significantly above Dec–Apr lows ($10,300–$11,200). Dealers who front-load acquisition in Q4 and Q1 and retail in summer capture a 2–4% seasonal swing — approximately $50k/year for a 100-unit operation.

Insight 7: Ontario, Tennessee, and Colorado are geographic premium markets.

These states show above-average price-to-MMR ratios. Dealers in low-price markets (the South and Midwest) can source vehicles regionally and transport to these markets for arbitrage.

Insight 8: Luxury and non-luxury pricing are entirely separate problems.

The Mann-Whitney test confirms the distributions are fundamentally different. A single model or pricing rule cannot serve both. Luxury segment pricing needs brand-specific benchmarks.

Insight 9: The Budget Fleet segment (Segment 1) is a volume game, not a margin game.

These $4k–$9k, high-mileage vehicles move in high volumes but carry thin margins. Operational efficiency (fast reconditioning, quick turn) is the only lever available in this segment.

Insight 10: Price deviations above 10% from MMR should trigger a manual review.

The deviation distribution is tight — most transactions are within ±5% of MMR. An outlier beyond ±10% almost always reflects either a data error, damage not captured in condition score, or an unusual market event.

11. Recommendations

The following five recommendations are prioritised by expected business impact, ordered from highest to lowest ROI. Each maps directly to an analytical insight and includes a feasibility note.

[CRITICAL] R1. Adopt MMR-Anchored Bid Pricing

Grounded In: Insight 1 — MMR explains 94% of selling price variance.

Recommendation: Mandate that all acquisition bids begin with the current MMR value as the baseline. Apply a standardised adjustment formula: Bid = MMR × (1 + mileage_adj + condition_adj + make_adj), where adjustment factors are derived from this project's regression model.

Expected Impact: Estimated reduction in acquisition overpayment of 1.5–2.5%, saving $200–$340 per vehicle. For a dealer acquiring 100 vehicles/month: $20k–$34k/month in additional gross profit.

Feasibility: Low implementation cost — update the buyer's bid worksheet template. Training required for acquisition staff. Full deployment within 30 days.

[HIGH] R2. Prioritise Low-Mileage Vehicle Sourcing

Grounded In: Insight 5 — 0–30k mile vehicles consistently sell above MMR.

Recommendation: Set a sourcing target of at least 30% of monthly acquisitions from the 0–30k odometer bucket. Specifically target off-lease vehicles from captive finance auctions, which typically have well-maintained condition profiles and mileage under 40,000.

Expected Impact: Above-MMR realisation in this bucket adds approximately $800–$1,200 per vehicle to retail gross profit. Shifting 30 units/month to this bucket: +$24k–$36k/month.

Feasibility: Requires relationship development with fleet/lease off-ramp sources. Higher acquisition cost partially offsets the margin gain — net analysis recommended per market.

[HIGH] R3. Implement Seasonal Acquisition Timing

Grounded In: Insight 6 — Mid-year prices (Jun–Jul) peak; Dec–Apr is the acquisition trough.

Recommendation: Front-load acquisition activity in Q4 (Oct–Dec) and Q1 (Jan–Mar) when auction prices are at seasonal lows, and retail aggressively through Q2–Q3 at peak prices. Model seasonal timing explicitly in inventory planning systems.

Expected Impact: A 2% average price reduction on 60 units acquired during the trough period saves $16,200/month over a 3-month acquisition window — approximately $50k/year.

Feasibility: No system change required — operational policy change only. Requires storage capacity for early-acquired inventory.

[MEDIUM] R4. Deploy Geographic Arbitrage Strategy

Grounded In: Insight 7 — Ontario, Tennessee, and Colorado show above-average price realisation.

Recommendation: For dealers in high-supply, low-price markets (e.g., Texas, Ohio, Indiana), identify and acquire vehicles that are in high demand in premium-price geographies. Transport economics (typically $600–$1,200 per vehicle) must be netted against the expected premium at the destination market.

Expected Impact: Where the geographic premium exceeds transport cost by 2× or more, net gain per vehicle of $500–$900. Even 10 vehicles/month on this strategy adds $60k–$108k/year.

Feasibility: Requires transport logistics partner and destination-market retail connection. Works best for premium late-model vehicles (Segment 3) with cleaner economics.

[MEDIUM] R5. Automate Condition-Based Pricing Adjustments

Grounded In: Insight 3 — Condition scores drive $2,000–$3,500 in price variation.

Recommendation: Integrate condition tier data into the dealership DMS (dealer management system) so that market price estimates are automatically adjusted when condition scores are entered. Alert buyers when their bid appears inconsistent with the vehicle's condition score relative to the segment average.

Expected Impact: Prevents systematic under-bidding on Excellent vehicles and over-bidding on Poor/Fair vehicles. Estimated margin improvement: $300–$500 per affected vehicle.

Feasibility: Requires DMS integration (API or CSV feed). Medium implementation complexity. Most modern DMS platforms support custom pricing rules.

12. Impact Estimation

The following table summarises the estimated financial impact of the five recommendations for a representative mid-size dealer group (100 vehicle acquisitions per month, 3 rooftops). All figures are approximations based on industry benchmarks and this project's analytical findings. The logic is sound; exact figures will vary by operator size, market, and execution quality.

Recommendation	Impact Type	Calculation Basis	Monthly Est.	Annual Est.
R1 — MMR-Anchored Bids	Cost Savings	100 units × $270 avg saving	$27,000 / month	$324,000 / year
R2 — Low-Mileage Sourcing	Revenue / Margin	30 units × $1,000 avg premium	$30,000 / month	$360,000 / year
R3 — Seasonal Timing	Cost Savings	60 units × 2% price reduction × $13,500	$16,200 / quarter	$50,000 / year
R4 — Geographic Arbitrage	Revenue Improvement	10 units × $700 net gain after transport	$7,000 / month	$84,000 / year
R5 — Condition Automation	Efficiency / Margin	~20 affected units × $400 correction	$8,000 / month	$96,000 / year
TOTAL (combined)	—	—	**~$88,200 / month**	**~$914,000 / year**

Why act now? The used-vehicle market is increasingly data-driven. Digital auction platforms, automated valuation models (AVMs), and AI-assisted bidding tools are already being deployed by large dealer groups and captive finance companies. Dealers who do not adopt data-backed pricing methodologies within the next 12–18 months will face a structural disadvantage against competitors who do. The tools exist, the data exists, and the ROI is quantified — the only variable is the decision to act.

13. Limitations

Data Constraints

- The dataset covers only 2014–2015. Market conditions, consumer preferences, fuel economy regulations, and EV adoption have changed substantially since then. Findings should be validated against current data before operational deployment.
- MMR values in the dataset represent the benchmark at time of sale — they cannot be backcasted or forward-projected without access to historical MMR time-series data, which is proprietary to Cox Automotive.
- Seller field quality is inconsistent. We cannot reliably identify whether specific auction houses or seller types (dealer vs fleet vs repo) systematically drive price deviations.
- The dataset does not include retail selling prices (consumer-facing). All insights are applicable to the wholesale auction market only.

Modelling Assumptions

- The multiple linear regression model assumes a linear relationship between predictors and selling price. In reality, depreciation is non-linear (especially in years 0–5 and above year 15). A log-transformed model or gradient boosting approach would likely improve accuracy.
- K-Means segmentation assumes spherical clusters of similar density. Actual vehicle segments may be non-spherical — hierarchical or DBSCAN clustering may reveal different segment structures.
- Hypothesis test results (all highly significant at $p < 0.001$) are partly a function of sample size ($n > 500{,}000$). Statistical significance does not automatically imply practical significance. Effect sizes should be interpreted alongside p-values.

What Cannot Be Concluded

- We cannot conclude that a specific vehicle will sell above or below MMR at a specific future auction — only that certain attributes are associated with above/below performance on average.
- We cannot draw conclusions about the retail (consumer) market, dealership profitability, or any dynamics beyond the wholesale auction floor.

14. Future Scope

Additional Analysis

- Expand the dataset to include 2016–2024 to capture post-pandemic volatility, EV pricing dynamics, and the impact of interest rate cycles on wholesale auction prices.
- Develop a non-linear price prediction model (XGBoost or Random Forest) that can outperform the linear regression, particularly for luxury and classic vehicle segments where depreciation curves are non-monotonic.
- Add make-model level fixed effects to the regression model to capture brand-specific systematic deviations from MMR that were observed in Figure 10.
- Perform survival analysis to model 'time-to-sale' at auction — not just what a vehicle sells for, but how many auction cycles it takes to clear. This is critical for inventory velocity planning.

New Data Sources

- Retail transaction data (consumer-facing prices from AutoTrader, Cars.com, or CarGurus) to build a wholesale-to-retail margin model.
- Vehicle condition inspection photos (for computer vision–based condition scoring automation to reduce subjective grader variance).
- Macroeconomic indicators (fuel prices, unemployment, consumer confidence) to model macro-driven auction price cycles.

Technology Roadmap

- Build a real-time pricing API that accepts vehicle attributes and returns a predicted price range with confidence interval, integrated into the buyer's mobile device at auction.
- Deploy a live Tableau dashboard connected to a real-time data feed for ongoing market monitoring, replacing the static 2014–2015 analysis.
- Explore conversational AI interfaces (natural language query on vehicle pricing data) for non-technical auction buyers who need price guidance without running SQL or Python.

15. Conclusion

The used-vehicle pricing challenge is, at its core, an information asymmetry problem: the data to make accurate pricing decisions exists, but it has historically not been organised, cleaned, and modelled in a way that is accessible to the buyers and managers who need it most. This project — **Judging a Car by its Cover** — directly addressed that gap.

Working with 548,486 real wholesale auction transactions, we built a complete analytics pipeline — from raw data ingestion through ETL cleaning, feature engineering, exploratory analysis, statistical modelling, customer segmentation, and interactive Tableau dashboards. The pipeline is reproducible, documented in version-controlled Jupyter notebooks, and designed to be re-run as new data becomes available.

The central insight is clear and actionable: **MMR is the strongest predictor of selling price (r = 0.97), but the 1.18% average discount and the wide ±50% deviation range mean that knowing the benchmark is necessary but not sufficient.** Mileage, vehicle age, condition score, transmission type, geographic market, and seasonal timing all move the needle — and we have quantified exactly how much each one matters.

Implementing the five recommendations in this report — anchoring bids to MMR with attribute-based adjustments, prioritising low-mileage sourcing, timing acquisitions seasonally, executing geographic arbitrage where the premium exceeds transport cost, and automating condition-based pricing alerts — is estimated to generate approximately **$914,000 in incremental annual gross profit** for a representative mid-size dealer group. The methodology is sound, the data is real, and the tools are ready.

16. Appendix

A. Full Data Dictionary

Column	Type	Description
year	int	Vehicle manufacture year (1990–2015)
make	str	Vehicle brand / manufacturer
model	str	Vehicle model name
trim	str	Trim level / variant
body	str	Body style (Sedan, SUV, Pickup, etc.)
transmission	str	automatic or manual
vin	str	17-character Vehicle Identification Number
state	str	US state code (2-letter) or Canadian province
condition	float	Condition score 1–49 (higher = better)
odometer	float	Odometer reading at time of sale (miles)
color	str	Exterior paint colour
interior	str	Interior trim colour
seller	str	Selling entity name
mmr	float	Manheim Market Report benchmark price (USD)
sellingprice	float	Actual transaction price (USD)
saledate	str	Sale timestamp (parsed to datetime)
sale_year	int	Calendar year from saledate
sale_month	int	Month number from saledate (1–12)
sale_month_name	str	Full month name
sale_dow	int	Day of week (0=Monday)
sale_dow_name	str	Day of week name
vehicle_age	int	sale_year minus year
price_deviation	float	sellingprice minus mmr
price_deviation_pct	float	(sellingprice – mmr) / mmr × 100
price_realization_rate	float	sellingprice / mmr × 100
sold_above_mmr	int	1 if sellingprice > mmr, else 0

price_per_age_year	float	sellingprice / vehicle_age
odometer_bucket	str	0-30k / 30-60k / 60-100k / 100-150k / 150k+
condition_tier	str	Poor / Fair / Good / Excellent (quartile-based)

B. Key Python Methodology Excerpts

```
# Saledate parsing — handles non-standard GMT format
import re
def parse_saledate(s):
s_clean = re.sub(r'\s+GMT.*', '', str(s)).strip()
return pd.to_datetime(s_clean, format='%a %b %d %Y %H:%M:%S')
```

```
# IQR-based winsorisation
def winsorise(df, col, k=2.5):
Q1, Q3 = df[col].quantile([0.25, 0.75])
IQR = Q3 - Q1
df[col] = df[col].clip(Q1 - k*IQR, Q3 + k*IQR)
return df
```

C. Cleaning Summary Log

Cleaning Step	Detail
Raw dataset loaded	558,837 rows × 16 columns
Rows with year < 1990 removed	98 rows dropped
Rows with sellingprice <= 0 removed	64 rows dropped
Transmission imputed (stratified mode)	65,330 values filled
Body imputed (make×model mode)	13,171 values filled
Condition imputed (make×body median)	11,856 values filled
Rows with make null dropped	10,287 rows dropped (after imputation not feasible)
sellingprice winsorised (2.5×IQR)	4,792 records capped
mmr winsorised (2.5×IQR)	5,237 records capped
odometer winsorised (2.5×IQR)	1,407 records capped
13 features engineered	vehicle_age, price_deviation, condition_tier, etc.
Final clean dataset	548,486 rows × 29 columns

17. Contribution Matrix

This section documents each team member's contribution across all project phases. All team members contributed across phases. Claims are verifiable through GitHub Insights, PR history, and committed notebook files.

Team Member (Enrolment No.)	Dataset & Sourcing	ETL & Cleaning	EDA & Analysis	Statistical Analysis	Tableau Dashboard	Report Writing	PPT
Abhinav Bajpai *(2401010017)* *Project Lead*	✓	✓			✓	✓	✓
Kushagra Bhardwaj *(2401010247)* *ETL Lead*	✓	✓				✓	
Rishiwant K. Maurya *(2401020102)* *Analysis Lead*			✓	✓		✓	
Pranav Singh *(2401010333)* *Visualization Lead*	✓		✓		✓	✓	
Samriddhi Shah *(2401010411)* *Strategy Lead*	✓		✓			✓	✓
Drishti Jha *(2401010162)* *PPT & Quality Lead*			✓			✓	✓

Declaration: We confirm that the above contribution details are accurate and verifiable through GitHub Insights, PR history, and submitted artifacts.